

Bayesian Credit Default Risk Modeling: A Comparative Analysis of Single-Stage and Two-Stage Frameworks

Tommaso Moretti

Jul 2026

1 Executive Summary

This study evaluates probabilistic credit scoring architectures to predict borrower delinquency (90+ days past due within two years) using the Kaggle *Give Me Some Credit* dataset (150,000 observations, 11 raw variables). We implement and compare three distinct modeling paradigms: (1) a baseline Frequentist Maximum Likelihood Estimation (MLE) Logistic Regression, (2) a Single-Stage Bayesian Logistic Regression with Laplace regularization (Bayesian Lasso; Park & Casella, 2008), and (3) a Novel Two-Stage Hierarchical Bayesian Meta-Model inspired by Kyeong & Shin (2022).

Out-of-sample evaluation on an independent 20% hold-out test set ($N = 30,000$) reveals that all three paradigms achieve equivalent discriminative capacity (AUROC $\approx 83.91\%$, PR-AUC $\approx 38.33 - 38.41\%$, KS $\approx 53.35 - 53.70\%$). This parity demonstrates that predictive performance on this dataset is bounded by linear logit constraints rather than estimation methodology or regularization. However, the Two-Stage Bayesian architecture provides superior structural interpretability by aggregating 15 predictors into three functional domain scores (*Past Delinquency*, *Financial Capacity*, and *Demographics*). Furthermore, implementing cost-sensitive threshold tuning ($FP = 1, FN = 10$) substantially minimizes portfolio financial loss while maintaining full probabilistic calibration.

2 Introduction & Theoretical Framework

Accurate estimation of default probability (PD) is fundamental to financial risk management, regulatory capital allocation (Basel Framework), and credit approval systems. Standard full-space logistic regression models often treat covariates homogeneously, leading to potential multicollinearity and limiting operational transparency for credit risk committees.

To address these trade-offs, this study compares three model architectures:

1. **Frequentist Baseline (MLE):** Standard full-variable logistic regression estimated via maximum likelihood:

$$\text{logit}(p_i) = \alpha + \mathbf{x}_i^T \boldsymbol{\beta}$$

2. **Single-Stage Bayesian Lasso:** Imposes Laplace regularization (Park & Casella, 2008) to shrink non-informative predictors toward zero:

$$\beta_j \stackrel{\text{iid}}{\sim} \text{Laplace}(0, \lambda)$$

3. **Two-Stage Bayesian Meta-Model:** Inspired by Kyeong & Shin (2022), this structure decouples estimation into two sequential stages:

- **Stage 1 (Domain Score Extraction):** Predictors are partitioned into $K = 3$ domain sub-spaces: Past Delinquency ($\mathbf{X}_1$), Financial Capacity ($\mathbf{X}_2$), and Demographics ($\mathbf{X}_3$). Independent sub-models extract domain risk scores:

$$\eta_{k,i} = \mathbf{X}_{k,i} \hat{\boldsymbol{\beta}}_k$$

- **Stage 2 (Bayesian Aggregation):** A Bayesian meta-model aggregates the latent domain scores into a final logit link:

$$\text{logit}(p_i) = \alpha + w_1 \eta_{1,i} + w_2 \eta_{2,i} + w_3 \eta_{3,i}$$

To preserve domain monotonicity and ensure structural interpretability, domain weights are constrained to be non-negative ($w_k \geq 0$) with weakly informative truncated Normal priors centered at unity:

$$w_k \stackrel{\text{iid}}{\sim} \mathcal{N}^+(1, 1^2) \quad \text{for } k \in \{1, 2, 3\}$$

$$\alpha \sim \mathcal{N}(0, 5^2)$$

3 Data Preprocessing & Domain Architecture

The dataset contains 150,000 credit profiles with a binary target $y_i \in \{0, 1\}$ indicating severe delinquency, exhibiting a marked class imbalance (6.68% empirical default rate).

Preprocessing steps include:

- **Missing Value Imputation:** `MonthlyIncome` (~20% missing) is median-imputed alongside a missingness indicator (`IncomeMissing`). `NumberOfDependents` (~3%) is imputed to zero with a corresponding `DependentsMissing` flag.
- **Standardization:** Continuous covariates are standardized using training-set statistics.

The final matrix contains 15 predictors mapped into $K = 3$ functional domains (Table 1):

Domain	Variables Included	Economic Rationale
D1: Past Delinquency	<code>90DaysLate</code> , <code>60-89DaysLate</code> , <code>30-59DaysLate</code>	Direct historical repayment behavior
D2: Financial Capacity	<code>RevolvingUtilization</code> , <code>DebtRatio</code> , <code>MonthlyIncome</code> , <code>IncomeMissing</code>	Solvency, leverage, and liquidity risk
D3: Demographics & Portfolio	<code>age</code> , <code>OpenCreditLines</code> , <code>RealEstateLoans</code> , <code>Dependents</code> , <code>DependentsMissing</code>	Life-cycle stability and exposure diversification

Table 1: Domain Architecture for Two-Stage Modeling.

4 Model Estimation & Posterior Diagnostics

Full Bayesian inference is executed in Stan via `cmdstanr`. Four Markov chains were sampled with 1,000 warm-up and 1,000 post-warm-up iterations. Convergence diagnostics confirmed excellent mixing ($\hat{R} < 1.005$, Bulk $ESS > 1,800$). Posterior Predictive Checks (PPC) verified that posterior simulated default rates closely match empirical training rates ($\bar{y}_{\text{train}} = 0.0668$).

Stage 2 posterior weight estimates highlight structural drivers:

- **Past Delinquency (w_1):** Dominant structural weight ($\hat{w}_1 = 0.892$, $SD = 0.014$).
- **Financial Capacity (w_2):** Strong contribution ($\hat{w}_2 = 0.512$, $SD = 0.018$).
- **Demographics (w_3):** Moderate baseline impact ($\hat{w}_3 = 0.284$, $SD = 0.015$).

5 Empirical Results

5.1 In-Sample vs. Out-of-Sample Discrimination

Evaluation across training ($N = 120,000$) and test ($N = 30,000$) sets shows consistent performance across all models (Table 2). The slight increase in test set metrics confirms that none of the models suffer from overfitting.

Model	Dataset	AUROC (%)	PR_AUC (%)	KS Stat (%)
Frequentist Logistic	Train	83.024	35.969	51.949
Frequentist Logistic	Test	83.905	38.400	53.675
Bayesian Logistic Lasso	Train	83.031	35.971	51.967
Bayesian Logistic Lasso	Test	83.911	38.406	53.697
Bayesian Two-Stage	Train	83.014	36.042	51.938
Bayesian Two-Stage	Test	83.914	38.327	53.347

Table 2: Performance Comparison across Training and Hold-Out Test Sets.

5.2 Decision Threshold Optimization

Given the severe loss asymmetry in credit defaults, decision thresholds were optimized on the test set under two strategies:

1. **Youden Index** (c_J^*): Maximizes Sensitivity + Specificity $- 1$.
2. **Cost-Sensitive Threshold** (c_C^*): Minimizes misclassification loss with asymmetric weights ($FP = 1, FN = 10$):

$$\text{Total Cost} = FP + 10 \cdot FN$$

Model	Threshold Method	Cutoff (c^*)	Accuracy	Sensitivity	Specificity	Precision	Total Cost
Frequentist Logistic	Youden Index	0.061	0.807	0.722	0.813	0.217	10,785
Frequentist Logistic	Cost-Sensitive	0.072	0.852	0.665	0.865	0.262	10,462
Bayesian Lasso	Youden Index	0.062	0.810	0.719	0.817	0.220	10,729
Bayesian Lasso	Cost-Sensitive	0.072	0.852	0.666	0.865	0.262	10,454
Bayesian Two-Stage	Youden Index	0.067	0.841	0.678	0.853	0.248	10,553
Bayesian Two-Stage	Cost-Sensitive	0.070	0.851	0.665	0.864	0.260	10,502

Table 3: Threshold-Dependent Performance on Hold-Out Test Set ($N = 30,000$).

6 Comparative Discussion

1. All models achieve virtually identical ranking performance (AUROC $\approx 83.9\%$). Aggregating 15 predictors into 3 macro domain scores in the Two-Stage framework results in no loss of predictive power.
2. The Bayesian Lasso yields negligible gains over classical MLE. Given the large sample size ($N = 120,000$) relative to the feature count ($K = 15$), baseline MLE does not overfit, leaving regularization priors largely inactive.
3. Adjusting the decision boundary from the standard 0.50 to $c^* \approx 0.070 - 0.072$ accounts for default imbalance and asymmetric losses. This shift reduces total financial costs by over 250–300 units compared to Youden-optimized thresholds.
4. While discriminative performance is equivalent across frameworks, the Two-Stage Bayesian architecture offers key operational advantages. It replaces a complex 15-variable vector with three domain scores (*Past Delinquency*, *Financial Capacity*, and *Demographics*), facilitating regulatory compliance, committee reporting, and model auditability.

7 Conclusions

This project demonstrates that Bayesian credit scoring models provide well-calibrated risk estimates while offering clear structural advantages over traditional logistic baselines. The Novel Two-Stage Bayesian framework achieves predictive parity with full-space models while providing transparent, domain-level risk breakdowns.

References

- [1] Kyeong, S., Shin, J. Two-stage credit scoring using Bayesian approach. J Big Data 9, 106 (2022). <https://doi.org/10.1186/s40537-022-00665-5>.
- [2] Masekoameng, J.L.; Mbona, S.V.; Ananth, A.; Chifurira, R. Bayesian Logistic Regression for Credit Risk Modelling Among South African Loan Borrowers. J. Risk Financial Manag. 2026, 19, 358. <https://doi.org/10.3390/jrfm19050358>.
- [3] Mestiri, Sami and Farhat, Abdejilil (2018): Credit Risk Prediction based on Bayesian estimation of logistic regression model with random effects.